



Applying research in stream processing for fraud detection

Pedro Cardoso
pedro.cardoso@feedzai.com

2019-10-25

whoami






whoami



Graduated from FEUP, MSc in computer science 2016

Obligatory tidbits:




- Love sports overall but  is **the** thing!
- Movies & live music on the streets.
-  
- Generally having a good time

whoami



Graduated from FEUP, MSc in computer science 2016

Obligatory tidbits:

- Love sports overall but  is **the** thing!
- Movies & live music on the streets.
-  
- Generally having a good time

@ **Feedzai** since **2017**.


Officially a **data engineer** in the **research team**.

Social networks:

- Github: <https://github.com/pedro93>
- Twitter: <https://twitter.com/pedro93>
- LinkedIn: <https://www.linkedin.com/in/pedro-silva-b7968733/>

So...

Stream Processing




Definition

*Processing an unlimited stream of data with finite resources
(Disk/RAM/CPU)*

Definition

*Processing an unlimited stream of data with finite resources
(Disk/RAM/CPU)*



There are some workarounds for that limitation...

- Scalar Functions -> *Sum/Count/Avg/Max/Min/Stddev*
- Temporal Queries -> Count of clicks in the past 5 minutes
- Sketches -> Count Min/Bloom Filters/Hyperloglog
- Infinite Pulse Responses -> Exponential averages
- Good Technology Choices
- Distributing Loads -> Partitions by group by keys
- Off-Loading State To Disk

Why

Stream processing is useful in fast, high data volume throughput systems:

High Frequency Trading

Fraud Detection

Clickstream analysis

Signal/Image/Video Processing

IoT Monitoring

Advertisement auctioning

Why

Stream processing is useful in fast, high data volume throughput systems:

High Frequency Trading

Fraud Detection

Clickstream analysis

Signal/Image/Video Processing

IoT Monitoring

Advertisement auctioning

And so much more...

The limit is your creativity



What solutions are out there



Spark
Streaming



HERON



Apache Flink



PULSAR



kafka



STORM




samza



Microsoft
Trill

There are others out there...

What solutions are out there



What is the best for you?

Depends, what are you looking for ?

What is the best for you?

Depends, what are you looking for ?

- Per-event accuracy or are you ok with micro batching?
- Fault Tolerant
- Low latency
- Adaptable
- Large Community
- Stateful vs Stateless
- High Throughput
- Scalable
- Battle tested

What is the best for you?

Depends, what are you looking for ?

- Per-event accuracy or are you ok with micro batching?
- Fault Tolerant
- Low latency
- Adaptable
- Large Community
- Stateful vs Stateless
- High Throughput
- Scalable
- Battle tested


Selecting one that has all these characteristics for your **specific use-case** is **HARD!**

@ Feedzai

@ Feedzai

PKernel

Streams to the rescue!



Profiles are temporal aggregations over some set of fields:


- Average spending in the last week.
- Distance to average location.
- Number of transactions in the last 10 minutes, ...

Very very useful as features for fraud-detecting ML Models & Rules

Lightweight Profiles

A research project using Exponential
Moving Averages

Context/Motivation



Context/Motivation

Naïve Sliding window profiles require:


- Overhead of storing events in the window (memory cost).
- Overhead to update/expire events in/out of the window.

Profiles used in projects: counts, sums, averages, count distincts, etc...


Alternatives to sliding windows :

- **Exponential Moving Average (EMA):** for counts, sums, averages, std dev.
- HyperLogLog (approximate counts): for count distincts.


Sliding window VS EMA



Sliding window VS EMA



Sliding window VS EMA



$$\begin{aligned} EMA(t_0) &= \sum_{i=0}^{+\infty} z_i e^{-\frac{t_0-t_i}{\tau}} \\ &= EMA(t_1) e^{-\frac{t_0-t_1}{\tau}} + z_0 \end{aligned}$$

Theoretically huge
memory savings!

In real-life

An EMA is everlasting, its value is never 0.

Not practical to store all real-time profiles since the beginning of time.

In real-life

An EMA is everlasting, its value is never 0.

Not practical to store all real-time profiles since the beginning of time.

Solution :

Key time-to-live: How long an EMA's "state" and its group by key should live.

In-memory data store of group-by key \Leftrightarrow metric state

*There goes our
memory savings?*


Overhead to update/expire events in/out of the window is now $O(1)$!

So did it work???

Latency Improvements

200 TPS
6.4M event real dataset
150 real-time profiles
over ~1.4M distinct values

Latency Improvements



200 TPS
6.4M event real dataset
150 real-time profiles
over ~1.4M distinct values

EMAs are **~2x to ~10x**
faster to compute than
Standard Operators

Memory Improvements


Actual values depend on configuration (key time-to-live factor).

EMAs can save **around 50%** memory in real-world scenarios.


Memory costs for profiles no longer depend on the number of events in the window.

LW-Conclusion

LW-Conclusion



WIPs



What are I'm working on right now...


Railgun

Redesigning internal engine to:

- Simultaneously handle massively large (> 1 year) & small (1 second) real-time windows in a transparent way.
- Compute accurate count-distincts over year-long-windows with millisecond latencies in real-time.

Lightweight Data Monitoring

Using approximate aggregations to detect shifts in data patterns in real-time with small computing resources.



Conclusion

There are a lot of streaming engines out there, but sometimes you need something custom.


Very few engines are distributed or do true event-by-event streaming, instead they micro-batch.

Traditional approaches to sliding windows are prone to bursting issues and are memory-intensive.

EMA-based features are lighter and do not lower performance of our ML models

In spite of decades of research, there is plenty of work to be done.

THANK YOU





Q&A